

# Homework Assignment #1

Jay Zhang

Due September 25th, 2024 at 11:59pm Pacific time

**IMPORTANT NOTE:** As indicated in the slides of Lecture 1 (pgs. 36-38) and the syllabus (pgs. 6-8), please list any resources outside of the course materials that you find helpful in completing the assignment (e.g. peers you discuss with, materials from different classes, blog posts, AI Tools, etc.). Please also be mindful of all policies in the syllabus concerning academic integrity and the use of AI Tools, including that you need to write your own solutions individually.

## Problem 1: (30 points)

Consider the simple linear regression model without an intercept:

$$Y = \beta X + \epsilon \quad (1)$$

Here there is a single feature  $X \in \mathbb{R}$  and a corresponding coefficient  $\beta \in \mathbb{R}$ . Note that the above model does not include an intercept term, i.e., it assumes that the intercept is zero. Suppose that we have collected data  $(x_1, y_1), \dots, (x_n, y_n)$ , where each  $x_i$  is an observed scalar value of the feature  $X$  and each  $y_i$  is a corresponding observed scalar value of the dependent variable  $Y$ .

Please answer the following:

- a) (5 points) Following the principle of minimizing the RSS error (i.e., ordinary least squares), derive a closed form solution for  $\hat{\beta}$ , the estimator of  $\beta$ , in terms of the data that we have collected.

model  $\rightarrow$  
$$Y = \beta X + \epsilon$$

$Y$  is labeled "dependent var" (with a line pointing to it).  
 $\beta$  is labeled "coefficient" (with a line pointing to it).  
 $X$  is labeled "independent var" (with a line pointing to it).  
 $\epsilon$  is labeled "error term" (with a line pointing to it).

Residual Sum of Squares (RSS) :

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$\downarrow$

$$\hat{y}_i = \beta x_i \rightarrow \text{so } RSS = \sum_{i=1}^n (y_i - \beta x_i)^2$$

minimizing KSS:

steps: ① take derivative of KSS with respect to  $\beta$

② set eq = 0

③ solve for  $\beta$

$$\begin{aligned} \textcircled{1} \quad \frac{d}{d\beta} KSS &= \frac{d}{d\beta} \sum_{i=1}^n (y_i - \beta x_i)^2 \\ &= 2 \sum_{i=1}^n (y_i - \beta x_i)(-x_i) = -2 \sum_{i=1}^n x_i (y_i - \beta x_i) \end{aligned}$$

$$\textcircled{2} \quad \text{set} = 0 \rightarrow -2 \sum_{i=1}^n x_i (y_i - \beta x_i) = 0$$

$$\textcircled{3} \quad \text{solve for } \beta \quad \sum_{i=1}^n x_i y_i - \beta \sum_{i=1}^n x_i^2 = 0$$

$$\sum_{i=1}^n x_i y_i = \beta \sum_{i=1}^n x_i^2$$

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2}$$

estimator of  $\beta$

Assume now that the collected data  $(x_1, y_1), \dots, (x_n, y_n)$  satisfies:

$$y_i = \beta x_i + \epsilon_i \text{ for all } i = 1, \dots, n,$$

where  $\beta \in \mathbb{R}$  is the true coefficient value and  $\epsilon_1, \dots, \epsilon_n$  are i.i.d. **normally distributed random variables** with **mean zero** and **variance  $\sigma^2$** . Here we assume that the  $x_i$  values (as well as the parameters  $\beta$  and  $\sigma$ ) are fixed and deterministic.

Under the above assumptions, please answer the following:

- b) (2 points) Explain why  $\hat{\beta}$  is a random variable.
- c) (6 points) Derive formulas for the expected value and variance of  $\hat{\beta}$ .
- d) (2 points) Argue that  $\hat{\beta}$  is normally distributed.

b) given  $y_i = \beta x_i + \epsilon_i$

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2} ; \text{ plug in } y_i$$

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i (\beta x_i + \epsilon_i)}{\sum_{i=1}^n x_i^2}$$

bc  $\hat{\beta}$  involves random variable  $\epsilon_i$ ,  
 which was said to be iid normal random var  
 $\Rightarrow \hat{\beta}$  depends on  $\epsilon_i$  (random var), making it also  
 a random var! ✓

$$\begin{aligned} c) \quad \hat{\beta} &= \frac{\sum_{i=1}^n x_i (\beta x_i + \epsilon_i)}{\sum_{i=1}^n x_i^2} = \frac{\sum_{i=1}^n \beta x_i^2 + \sum_{i=1}^n x_i \epsilon_i}{\sum_{i=1}^n x_i^2} \\ &= \beta + \frac{\sum_{i=1}^n x_i \epsilon_i}{\sum_{i=1}^n x_i^2} \end{aligned}$$

expected val of  $\hat{\beta}$

$$E[\hat{\beta}] = E\left[\beta + \frac{\sum_{i=1}^n X_i \varepsilon_i}{\sum_{i=1}^n X_i^2}\right]$$

linearity + expectation

$$E[X \cdot Y] = E[X] \cdot E[Y] \rightarrow E[\hat{\beta}] = \beta + \frac{1}{\sum_{i=1}^n X_i^2} E\left[\sum_{i=1}^n X_i \varepsilon_i\right]$$

/ constant

given: mean of  $\varepsilon_i = 0 \rightarrow E[\varepsilon_i] = 0$  for all  $i$

$$E\left[\sum_{i=1}^n X_i \varepsilon_i\right] = 0$$

$$\text{thus } E[\hat{\beta}] = \beta + 0 = \beta$$

variance of  $\hat{\beta}$ :

bc constant, can be ignored since  $\text{var} = 0$

$$\text{var}(cZ) = c^2 \text{var}(Z)$$

$$\text{var}(\hat{\beta}) = \text{var}\left(\beta + \frac{\sum_{i=1}^n X_i \varepsilon_i}{\sum_{i=1}^n X_i^2}\right) = \frac{1}{\left(\sum_{i=1}^n X_i^2\right)^2} \cdot \text{var}\left(\sum_{i=1}^n X_i \varepsilon_i\right)$$

$\varepsilon_i \Rightarrow \text{iid}, \text{var } \sigma^2$

$$\begin{aligned}\text{var}\left(\sum_{i=1}^n X_i \varepsilon_i\right) &= \sum_{i=1}^n X_i^2 \text{var}(\varepsilon_i) \\ &= \left(\sum_{i=1}^n X_i^2\right) \cdot \sigma^2\end{aligned}$$

$$\text{thus } \text{var}(\hat{\beta}) = \frac{\sigma^2}{\sum_{i=1}^n X_i^2}$$

a) argue that  $\hat{\beta}$  is normally distributed.

$$\hat{\beta} = \beta + \frac{\sum_{i=1}^n X_i \epsilon_i}{\sum_{i=1}^n X_i^2} \rightarrow \text{linear combo of } \epsilon_i$$

mean = 0, var  $\sigma^2$   
normal random var, iid  
 $\downarrow$   
 bec this holds, then the  
 linear combo of it also is normally distr

then  $\sum_{i=1}^n X_i \epsilon_i \Rightarrow$  normally distr with mean = 0  
 var =  $\sigma^2 \sum_{i=1}^n X_i^2$

thus  $\frac{\sum_{i=1}^n X_i \epsilon_i}{\sum_{i=1}^n X_i^2} \Rightarrow$  normally distr with mean = 0  
 var =  $\frac{\sigma^2}{\sum_{i=1}^n X_i^2}$  ✓

Let  $x_{n+1}$  be a newly observed feature value, with associated value of the dependent variable  $y_{n+1}$ , satisfying:

$$y_{n+1} = \beta x_{n+1} + \epsilon_{n+1},$$

where  $\epsilon_{n+1}$  is independent of  $\epsilon_1, \dots, \epsilon_n$  and follows the same distribution. Currently, we have observed  $x_{n+1}$  but have not yet observed  $y_{n+1}$ .

Please answer the following:

- e) (5 points) Assuming that  $\sigma$  is known but  $\beta$  is not, describe how to construct a 95% *confidence interval* for  $\beta x_{n+1}$ . This is a random interval, constructed from the observed data and the value of  $\sigma$ , such that the probability that the interval contains  $\beta x_{n+1}$  is 95%. This probability is calculated over the randomness associated with the observed data  $(x_1, y_1), \dots, (x_n, y_n)$ .
- f) (5 points) Assuming that both  $\sigma$  and  $\beta$  are known, describe how to construct a 95% *prediction interval* for  $y_{n+1}$ . This is an interval such that the probability that  $y_{n+1}$  lies in the interval is 95%.
- g) (5 points) Assume now that  $\sigma$  is known but  $\beta$  is not. Describe how to construct a 95% *prediction interval* for  $y_{n+1}$ . This is a random interval, constructed from the observed data and the value of  $\sigma$ , such that the probability that  $y_{n+1}$  lies in the interval is 95%. This probability is calculated over all randomness associated with  $(x_1, y_1), \dots, (x_n, y_n), (x_{n+1}, y_{n+1})$ .

**NOTE:** if you choose to use linear algebra techniques to address any part of this problem (which is not required), please simplify all answers in terms of summations instead of matrix/vector notation.

$$e) \quad \hat{\beta} = \frac{\sum_{i=1}^n X_i y_i}{\sum_{i=1}^n X_i^2} \quad SE(\hat{\beta}) = \frac{\sigma}{\sqrt{\sum_{i=1}^n X_i^2}} \quad \begin{array}{l} \sigma \text{ known} \\ \beta \text{ unknown} \end{array}$$

$$95\% \text{ confidence interval for } \hat{\beta}_{n+1} \Rightarrow \left( \hat{\beta}_{n+1} - 1.96 \cdot \frac{\sigma x_{n+1}}{\sqrt{\sum_{i=1}^n X_i^2}}, \hat{\beta}_{n+1} + 1.96 \cdot \frac{\sigma x_{n+1}}{\sqrt{\sum_{i=1}^n X_i^2}} \right)$$

f)  $\sigma$  known  $\beta$  known

$$\text{from } \hat{\beta} = \frac{\sum_{i=1}^n X_i y_i}{\sum_{i=1}^n X_i^2} \Rightarrow \hat{y}_{n+1} = \hat{\beta} X_{n+1}$$

given  $y_{n+1} = \beta X_{n+1} + \varepsilon_{n+1} \Rightarrow$  know also  $= \hat{y}_{n+1} + \varepsilon_{n+1}$  / Sub in for  $\beta X_{n+1}$

prediction interval  $\Rightarrow y_{n+1} = \hat{y}_{n+1} + \varepsilon_{n+1}$  where  $\varepsilon_{n+1} \sim N(0, \sigma^2)$

for a 95% prediction interval:  $\hat{y}_{n+1} \pm z_{\alpha/2} \cdot \sigma$

$$= (\hat{\beta} X_{n+1} - 1.96 \sigma, \hat{\beta} X_{n+1} + 1.96 \cdot \sigma)$$

g)  $\sigma$  known  $\beta$  unknown

estimate  $\hat{\beta} = \frac{\sum_{i=1}^n X_i y_i}{\sum_{i=1}^n X_i^2} \quad \hat{y}_{n+1} = \hat{\beta} X_{n+1}$

when constructing a prediction interval for  $y_{n+1}$ , need to consider both ① uncertainty of fitted model and ② randomness of new data pt

$\hookrightarrow$  overall uncertainty = sum of 2 vars

$$\Rightarrow \text{var}(y_{n+1}) = \text{var}(\hat{y}_{n+1}) + \sigma^2 = \frac{\sigma^2 X_{n+1}^2}{\sum_{i=1}^n X_i^2} + \sigma^2$$

prediction interval = range of vals that contain  $y_{n+1}$  with 95% confidence

$$\hat{\beta} X_{n+1} \pm z_{\alpha/2} \cdot \sqrt{\frac{\sigma^2 X_{n+1}^2}{\sum_{i=1}^n X_i^2} + \sigma^2}$$

$$\left[ \hat{\beta} X_{n+1} - 1.96 \sqrt{\frac{\sigma^2 X_{n+1}^2}{\sum_{i=1}^n X_i^2} + \sigma^2}, \hat{\beta} X_{n+1} + 1.96 \sqrt{\frac{\sigma^2 X_{n+1}^2}{\sum_{i=1}^n X_i^2} + \sigma^2} \right]$$

# 142AHW1

September 26, 2024

## 0.1 IEOR 142A: Introduction to Machine Learning and Data Analytics I, Fall 2024

### 0.2 Homework Assignment #1

#### 0.2.1 Problem 2: Forecasting Honda Civic Sales

```
[103]: from pandas.plotting import scatter_matrix
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import math
import random
import statsmodels.api as sm
random.seed(30)
```

```
[104]: #load csv data
civic = pd.read_csv("civiccar.csv")
civic
```

```
[104]:
```

	MonthNumeric	MonthFactor	Year	CivicSales	Unemployment	CivicQueries	\
0	1	January	2014	21824	6.6	66	
1	2	February	2014	21575	6.7	69	
2	3	March	2014	27697	6.7	72	
3	4	April	2014	27611	6.2	69	
4	5	May	2014	36089	6.3	69	
..	...	...	...	...	...	...	
122	3	March	2024	5664	3.8	87	
123	4	April	2024	5348	3.9	83	
124	5	May	2024	6700	4.0	88	
125	6	June	2024	5935	4.1	92	
126	7	July	2024	5755	4.3	92	

	CPIA11	CPIEnergy	MilesTraveled	interest
0	235.288	250.340	246531	NaN
1	235.547	249.925	249499	NaN
2	236.028	249.961	251120	NaN
3	236.468	249.864	251959	NaN

```

4      236.918      249.213      252289      NaN
..      ...      ...      ...      ...
122    312.230      287.399      273352      8.877857
123    313.207      290.631      273430      17.200000
124    313.225      284.742      274175      8.877857
125    313.049      278.938      274160      8.877857
126    313.534      279.012      274273      17.080000

```

[127 rows x 10 columns]

- a) (25 points) Start by splitting the data into a training set and a testing set. The training set should contain all observations from 2014 through 2019. The testing set should have all observations from January 2020 through July 2024

```

[105]: #splitting data into training and test set- filtering years
civic_train = civic[civic['Year'] <= 2019]
civic_test = civic[civic['Year'] >= 2020]

```

```

[106]: civic.info() #looking at different features

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 127 entries, 0 to 126
Data columns (total 10 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   MonthNumeric    127 non-null    int64
 1   MonthFactor     127 non-null    object
 2   Year            127 non-null    int64
 3   CivicSales      127 non-null    int64
 4   Unemployment    127 non-null    float64
 5   CivicQueries    127 non-null    int64
 6   CPIAll          127 non-null    float64
 7   CPIEnergy       127 non-null    float64
 8   MilesTraveled   127 non-null    int64
 9   interest        116 non-null    float64
dtypes: float64(4), int64(5), object(1)
memory usage: 10.1+ KB

```

```

[107]: features = ['CivicSales',
                  'Unemployment',
                  'CivicQueries',
                  'CPIEnergy',
                  'CPIAll',
                  'MilesTraveled']

```

```


[108]: corr = civic[features].corr()

```



```
# Plot clustermap to showcase correlation
sns.clustermap(corr, annot=True, cmap='coolwarm', figsize=(8, 8))
```

[108]: <seaborn.matrix.ClusterGrid at 0x7f420ec4a690>



## Model 1: Regular OLS

i) **Feature Selections:** Through the displayed diagram, we can finalize our selection of features correlating best to Civic Sales:

- **CPIAll:** CPI-All measures inflation by tracking price changes for a basket of goods and services commonly purchased by urban consumers, including car sales. It reflects cost-of-living changes and is a key indicator used to assess inflation in urban areas.

- **CivicQueries:** CivicQueries is included because search frequency and behavior are often reasonable indicators in gauging the interest of the population in specific products/services, which can be beneficial for our case in understanding Car Sales.
- **MilesTraveled:** MilesTraveled displays a potential increase in wants and desires. If more people are driving/traveling, this could lead to overuse of their cars, which could lead to them needing to buy a new one because of tears or malfunction issues.
- **Unemployment:** Unemployment was selected because economic fluctuations are strong indicators of consumer demand. With overburdening economic stress and without a surplus of money, this could lead to shifts in consuming behaviors/patterns– the overall population may not choose to spend their money on purchasing vehicles OR choosing cheaper alternatives.
- **CPIEnergy** was excluded from my final model because it is a subset of CPIAll and I did not believe it was relevant to include both. This variable could even be capable of creating collinearity issues with CPIAll, which could negatively impact our model. While comparing the accuracy of the model, it also just seemed overall better after we excluded CPIEnergy and kept the rest of the variables that seemed more linearly independent of one another.

**ii) Our Model:** Using OLS with these Civic Sales as our dependent variable and the 4 independent variables, we have the following equation:

$$Y = -33,180 - 149.39 \cdot X_{CPIA} + 366.19 \cdot X_{CQ} + 0.220 \cdot X_{MT} + 2535.01 \cdot X_u$$

**iii) Coefficient Interpretation:**

- **Intercept -33,180:** When all the other variables are 0, this would be the value we obtain in sales. Because the intercept is negative, it does not make logical/statistical sense on its own, as you cannot really incur negative sales. One reason for this could be due to the relative collinearity in the chosen features.
- **CPIAll -149.39:** A 1-unit increase in the overall CPIAll is associated with a decrease in Civic sales by 149 units. A rise in the CPIAll can lead to a decrease in Honda Civic sales by impacting consumer confidence, increasing the overall cost of vehicle ownership, and making financing less accessible.
- **CivicQueries 366.19:** A 1-unit increase in Civic-related online searches is associated with a 366 Civic sale increase, indicating a pretty positive/strong relationship between car sales and internet search queries since people likely do a lot of research on cars before they decide to purchase. This is logically reasonable, as we do see the skyrocket of online traffic correlating to real-world interactions and surges in purchases.
- **MilesTraveled 0.22:** For every 1-unit increase in miles traveled, Civic Sales are predicted to increase by 0.22 units. Compared to the other variables, this change is statistically insignificant, indicating that it likely has a minimal impact on Civic Sales overall.
- **Unemployment 2,535.01:** For every 1% increase in unemployment, Civic Sales are predicted to increase by 2,535 units. Surprised that the coefficient was actually positive, this could imply that during times when there is an increase in unemployment, people might opt to buy cheaper alternatives for cars, like the Honda Civic compared to expensive/luxury cars. If the

coefficient was negative, that would follow with my original thought, where people in general would be less inclined to buy cars in general, leading to a decrease in sales.

```
[109]: model1_features = ['CPIA11', 'CivicQueries',
                        'MilesTraveled', 'Unemployment']

X_train_1 = civic_train[model1_features]
X_train_1 = sm.add_constant(X_train_1)
y_train_1 = civic_train['CivicSales']
```

```
[110]: model1 = sm.OLS(y_train_1, X_train_1).fit()
print(model1.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  CivicSales    R-squared:                  0.412
Model:                            OLS        Adj. R-squared:              0.377
Method:                 Least Squares    F-statistic:                  11.75
Date:                Thu, 26 Sep 2024    Prob (F-statistic):          2.75e-07
Time:                      01:37:04    Log-Likelihood:              -688.83
No. Observations:                  72      AIC:                        1388.
Df Residuals:                      67      BIC:                        1399.
Df Model:                          4
Covariance Type:                  nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-
const                -3.318e+04    6.61e+04    -0.502    0.617    -1.65e+05
9.87e+04
CPIA11                -149.3855    141.679    -1.054    0.295    -432.178
133.407
CivicQueries           366.1866    63.883     5.732    0.000     238.675
493.698
MilesTraveled          0.2202     0.187     1.178    0.243     -0.153
0.594
Unemployment          2535.0126    1844.274     1.375    0.174    -1146.173
6216.198
=====
Omnibus:                 1.861    Durbin-Watson:              1.576
Prob(Omnibus):           0.394    Jarque-Bera (JB):           1.427
Skew:                   0.142    Prob(JB):                   0.490
Kurtosis:               2.372    Cond. No.                   4.13e+07
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.13e+07. This might indicate that there are strong multicollinearity or other numerical problems.

iv) **Results and Observations:** Our model's **R-squared** value is **0.412**, indicating that the model explains approximately **41.2%** of the variance seen within our Civic Sales. Thus we were able to capture around half of the factors influencing sales.

- The **F-statistic** of **11.75** with a p-value of **2.75e-07** shows that the model is statistically significant overall, meaning a statistically meaningful relationship was created with at least one of the independent variable features. The statistic overall tells us whether the predictors in our model are useful for predicting the outcome of Civic Sales.
- The **Adjusted R-squared** of **0.377** tells us how well our model fits the data and accounts for the number of predictors we have included.

Final Conclusion: While our model is pretty reasonable in its predictions, we could have performed better if we were not limited to the select features given to us that could be able to encompass more of what influences Civic Sales.

## Model 2: Seasonality

i) **Our Model:**  $\$CivicSales = \theta + 1 X + 2 X\{CQ\} + 3 X\{MT\} + 4 X\{CPIE\} + 5 X\{CPIA\} + \sum_{i=1}^{11} \{i+5\} X\{MFi\}$

**Coefficient Interpretation:** No Variable Selection Occurred in this problem.

- Intercept
- Unemployment
- CivicQueries
- MilesTraveled
- CPIEnergy
- CPIAll
- Month Factor: **new** independent variable in addition to all five of the variables we used at the start of Part (a). Seasonality is important in predicting demand and sales since demand for many products tends to be periodic in time. In regards to Civic Sales, people may buy cars more during winter since there are more discounts that occur during Black Friday or Christmas to incentivize consumers. Also, people may just be in a more spending habit/spirit due to festivities. Here, each dummy variable of Month Factor accounts for every month (Jan to Dec) in a year.

## ii) iii) Results and Observations

```
[111]: model2_features =  
        ['Unemployment', 'CivicQueries', 'CPIEnergy', 'CPIAll', 'MilesTraveled']
```

```

# One-hot encode 'MonthFactor' variable
X_train_months = pd.get_dummies(civic_train['MonthFactor'], drop_first=True)
X_train_2 = pd.concat([civic_train[model2_features], X_train_months], axis=1)

#boolean to integer
X_train_2 = X_train_2.astype(int)

#merge with original features
X_train_2 = sm.add_constant(X_train_2)
y_train_2 = civic_train['CivicSales']

```

```

[112]: model2 = sm.OLS(y_train_2, X_train_2).fit()
print(model2.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  CivicSales    R-squared:                  0.806
Model:                            OLS        Adj. R-squared:              0.749
Method:                 Least Squares    F-statistic:                  14.25
Date:                Thu, 26 Sep 2024    Prob (F-statistic):          3.86e-14
Time:                      01:37:04    Log-Likelihood:              -649.00
No. Observations:                72      AIC:                        1332.
Df Residuals:                    55      BIC:                        1371.
Df Model:                        16
Covariance Type:                nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-
const          -5.08e+04    3.11e+04    -1.632    0.108    -1.13e+05
1.16e+04
Unemployment    2942.7866    871.126     3.378    0.001    1197.010
4688.563
CivicQueries    258.0527    65.154     3.961    0.000    127.480
388.625
CPIEnergy       -8.0006    30.288    -0.264    0.793    -68.700
52.699
CPIA11          -117.6125    152.561    -0.771    0.444    -423.352
188.127
MilesTraveled    0.2922    0.148     1.977    0.053    -0.004
0.589
August          3417.9077    1377.606     2.481    0.016    657.124
6178.692
December        2446.2169    1574.916     1.553    0.126    -709.985

```

5602.418					
February	-3508.0232	1393.483	-2.517	0.015	-6300.626
-715.420					
January	-5124.7768	1374.966	-3.727	0.000	-7880.270
-2369.283					
July	1502.8602	1387.625	1.083	0.284	-1278.002
4283.723					
June	1167.4982	1400.212	0.834	0.408	-1638.589
3973.585					
March	1267.1779	1383.511	0.916	0.364	-1505.441
4039.797					
May	5269.6578	1395.259	3.777	0.000	2473.495
8065.820					
November	-940.8309	1464.543	-0.642	0.523	-3875.841
1994.179					
October	-1412.3330	1382.125	-1.022	0.311	-4182.174
1357.508					
September	-689.1123	1368.468	-0.504	0.617	-3431.584
2053.359					
=====					
Omnibus:	10.476	Durbin-Watson:	1.500		
Prob(Omnibus):	0.005	Jarque-Bera (JB):	12.013		
Skew:	0.680	Prob(JB):	0.00246		
Kurtosis:	4.468	Cond. No.	3.07e+07		
=====					

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.07e+07. This might indicate that there are strong multicollinearity or other numerical problems.

Our model's **R-squared** value is **0.806**, indicating that the model explains approximately **80.6%** of the variance seen within our Civic Sales. Comparing this to our first simple model, we were able to capture a lot more variance around the factors influencing our Civic Car Sales.

Due to the big jump in improvement from part A, we can attribute a lot of this to our new variable **Month Factor**. It seems like seasonality truly has an impact on the performance of our Car Sales. Based on the chart we have generated, we see an impact of increases in Car Sales during the months of August, December, and May, which could be explained as starting school/new job, during holidays/festivities/gift giving/discount season, and end-of-school-year celebrations. Looking back to our previous variables, we see that Unemployment, followed by CivicQueries, and CPIALL has stronger coefficients that impact our model, thereby impacting Civic Sales.

**Final Conclusion:** Month Factor was a good choice in adding to our model since it seemed to encompass/capture our variance a lot better than our first model was able to do with just the 4 features. Our higher R-squared value shows an overall increase in assessing the explanatory power of our regression model.

iv) **Another Way to Model Seasonality** We could potentially look into other models that are known to capture seasonality well. One example that is commonly used is **the Seasonal Autoregressive Integrated Moving Average (SARIMA)**, a powerful and versatile model for time series forecasting that effectively incorporates seasonal patterns while providing a solid framework for understanding and predicting data trends. If the data exhibits strong seasonal patterns, a model specifically designed to capture seasonality like SARIMA might outperform a more general model like the one we had. From what we observed in the previous model implementing Month Factor, it does seem like seasonality is strongly correlated to our model, thus I do believe using a model like SARIMA would be able to improve our performance.

### Model 3: Mixed Model

```
[113]: model3_features = ['Unemployment', 'CivicQueries', 'CPIAll', 'MilesTraveled']

# One-hot encode 'MonthFactor' and merge
X_train_with_month = pd.get_dummies(civic_train['MonthFactor'], drop_first=True)
X_train_3 = pd.concat([civic_train[model3_features], X_train_with_month],
    ↪axis=1)

X_train_3 = X_train_3.astype(int) #bool to int

X_train_3 = sm.add_constant(X_train_3)
y_train_3 = civic_train['CivicSales']

[114]: model3 = sm.OLS(y_train_3, X_train_3).fit()
print(model3.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  CivicSales      R-squared:                0.805
Model:                          OLS           Adj. R-squared:           0.753
Method:                        Least Squares   F-statistic:              15.45
Date:                          Thu, 26 Sep 2024 Prob (F-statistic):       9.89e-15
Time:                          01:37:04       Log-Likelihood:           -649.04
No. Observations:              72             AIC:                     1330.
Df Residuals:                  56             BIC:                     1367.
Df Model:                      15
Covariance Type:               nonrobust
=====
=
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-
const          -5.024e+04    3.08e+04    -1.631    0.109    -1.12e+05
1.15e+04
Unemployment    2845.4641     782.791     3.635    0.001    1277.345

```

4413.583					
CivicQueries	260.5953	63.902	4.078	0.000	132.584
388.607					
CPIA11	-150.2400	88.795	-1.692	0.096	-328.117
27.637					
MilesTraveled	0.3145	0.120	2.611	0.012	0.073
0.556					
August	3436.6776	1364.298	2.519	0.015	703.661
6169.694					
December	2485.6162	1554.760	1.599	0.116	-628.943
5600.175					
February	-3450.5977	1364.940	-2.528	0.014	-6184.902
-716.293					
January	-5070.8322	1348.374	-3.761	0.000	-7771.951
-2369.714					
July	1517.3142	1374.981	1.104	0.275	-1237.104
4271.733					
June	1207.2495	1380.490	0.875	0.386	-1558.205
3972.704					
March	1321.3146	1356.835	0.974	0.334	-1396.753
4039.382					
May	5279.3923	1383.140	3.817	0.000	2508.630
8050.154					
November	-863.6963	1423.166	-0.607	0.546	-3714.641
1987.248					
October	-1384.5645	1366.627	-1.013	0.315	-4122.248
1353.119					
September	-638.6332	1343.757	-0.475	0.636	-3330.503
2053.236					
=====					
Omnibus:	10.047	Durbin-Watson:	1.495		
Prob(Omnibus):	0.007	Jarque-Bera (JB):	11.366		
Skew:	0.659	Prob(JB):	0.00340		
Kurtosis:	4.432	Cond. No.	3.06e+07		
=====					

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.06e+07. This might indicate that there are strong multicollinearity or other numerical problems.

#### Results and Observations:

Our model's **R-squared** value is **0.805**, indicating that the model explains approximately **80.5%** of the variance seen within our Civic Sales. Comparing to the first 2 models we created, there was a very slight decrease in the value compared to our 2nd model.

Comparing to the **OSR Squared** value below (0.5105) to our R-squared value (0.805), we can



see that the OSR Squared value is smaller. When R-squared is greater than OSR Squared, that usually implies overfitting (aka the model fits the training data better than it generalizes to new data). We could infer that our model might be a little too complex and capture noise rather than the underlying pattern. To mitigate this, we might consider simplifying the model or using regularization techniques to avoid overfitting.

```
[115]: #Obtained function from AI as I could not find the one from class that works
        ↪with code
def calculate_osr2_from_ols(r_squared, n, df_model, df_residual):

    # Calculate total sum of squares (SS_total)
    ss_total = (n - 1) # Using a normalized approach (this assumes the
    ↪variance of Y is 1)

    # Calculate SS_regression and SS_residual
    ss_regression = r_squared * ss_total
    ss_residual = ss_total - ss_regression

    # Calculate OSR2
    osr2 = (ss_regression - ss_residual) / (ss_total + ss_residual)

    return osr2

#from our regression results
df_model = 15
df_residual = 56
r_squared = 0.805
n = 72

osr2_value = calculate_osr2_from_ols(r_squared, n, df_model, df_residual)
print(f"OSR2 Value: {osr2_value:.4f}")
```

OSR<sup>2</sup> Value: 0.5105

**Additional Variable to Implement** Another variable that could be related to Honda sales is **monthly interest rates on auto loans**. When interest rates are low, consumers are more likely to finance car purchases, which could lead to higher sales of Honda vehicles. Auto loan rates are a key factor in consumer decision-making when purchasing cars, so adding this variable to our regression model could provide insights into whether lower rates are driving higher Honda sales.

```
[116]: #load csv data
new = pd.read_csv("newfeature.csv")
new
```

```
[116]:      DATE  TERMCBAUTO048NS
0  1972-02-01          10.20
1  1972-03-01              .
2  1972-04-01              .
```

3	1972-05-01	9.96
4	1972-06-01	.
..	...	...
623	2024-01-01	.
624	2024-02-01	8.57
625	2024-03-01	.
626	2024-04-01	.
627	2024-05-01	8.65

[628 rows x 2 columns]

[117]: *#splitting data into training and test set- filtering years*

```
# Convert the 'date' column to datetime
new['DATE'] = pd.to_datetime(new['DATE'])
# Extract the year and create a new column for it
new['year'] = new['DATE'].dt.year

new.replace('.', pd.NA, inplace=True)

#Convert the column to a numeric type (this will handle non-numeric values like
↪ '.')
new['TERMCBAUTO48NS'] = pd.to_numeric(new['TERMCBAUTO48NS'], errors='coerce')

#Replace NaN values with the mean of the column
new['TERMCBAUTO48NS'].fillna(new['TERMCBAUTO48NS'].mean(), inplace=True)

new_filter= new[new['year'] > 1972]
```

/tmp/ipykernel\_357/2309341815.py:14: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
new['TERMCBAUTO48NS'].fillna(new['TERMCBAUTO48NS'].mean(), inplace=True)
```

[118]: *# Step 3: Select the specific column from the first DataFrame*  
*# Replace 'column\_name' with the actual name of the column you want to import*  
interest = new['TERMCBAUTO48NS']

```

# Step 4: Add the selected column to the second DataFrame
# Ensure the lengths match, or consider how to handle differing lengths
civic['interest'] = interest # Replace 'new_column_name' with your desired_
    ↳column name

# Step 5: Save the updated DataFrame back to a CSV file
civic.to_csv('civiccar.csv', index=False)

print(civic)

```

	MonthNumeric	MonthFactor	Year	CivicSales	Unemployment	CivicQueries	\
0	1	January	2014	21824	6.6	66	
1	2	February	2014	21575	6.7	69	
2	3	March	2014	27697	6.7	72	
3	4	April	2014	27611	6.2	69	
4	5	May	2014	36089	6.3	69	
..	...	...	...	...	...	...	
122	3	March	2024	5664	3.8	87	
123	4	April	2024	5348	3.9	83	
124	5	May	2024	6700	4.0	88	
125	6	June	2024	5935	4.1	92	
126	7	July	2024	5755	4.3	92	

	CPIAll	CPIEnergy	MilesTraveled	interest
0	235.288	250.340	246531	10.200000
1	235.547	249.925	249499	8.877857
2	236.028	249.961	251120	8.877857
3	236.468	249.864	251959	9.960000
4	236.918	249.213	252289	8.877857
..	...	...	...	...
122	312.230	287.399	273352	8.877857
123	313.207	290.631	273430	17.200000
124	313.225	284.742	274175	8.877857
125	313.049	278.938	274160	8.877857
126	313.534	279.012	274273	17.080000

[127 rows x 10 columns]

[119]: #splitting data into training and test set- filtering years

```

civic['interest'].fillna(civic['interest'].mean(), inplace=True)
civic_train = civic[civic['Year'] <= 2019]
civic_test = civic[civic['Year'] >= 2020]
model4_features =
    ↳['Unemployment', 'CivicQueries', 'CPIAll', 'MilesTraveled', 'interest']

X_train_with_month = pd.get_dummies(civic_train['MonthFactor'], drop_first=True)

```

```

X_train_4 = pd.concat([civic_train[model4_features], X_train_with_month],
↳axis=1)

X_train_4 = X_train_4.astype(int) #bool to int

X_train_4 = sm.add_constant(X_train_4)
y_train_4 = civic_train['CivicSales']

model4 = sm.OLS(y_train_4, X_train_4).fit()
print(model4.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          CivicSales    R-squared:                0.812
Model:                  OLS          Adj. R-squared:            0.757
Method:                 Least Squares  F-statistic:              14.83
Date:                  Thu, 26 Sep 2024  Prob (F-statistic):      1.68e-14
Time:                  01:37:04       Log-Likelihood:           -647.84
No. Observations:      72            AIC:                     1330.
Df Residuals:          55            BIC:                     1368.
Df Model:              16
Covariance Type:       nonrobust
=====

```

```

=====
=
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-
const      -5.012e+04   3.06e+04    -1.639    0.107   -1.11e+05
1.11e+04
Unemployment  3011.4380    786.184     3.830    0.000    1435.889
4586.987
CivicQueries  253.1797     63.642     3.978    0.000     125.639
380.720
CPIA11      -179.1112     90.603    -1.977    0.053    -360.685
2.462
MilesTraveled  0.3915      0.132     2.963    0.004      0.127
0.656
interest    -1346.4870    984.179    -1.368    0.177   -3318.827
625.853
August       680.7725    2427.017     0.280    0.780   -4183.078
5544.622
December    -531.7662    2691.534    -0.198    0.844   -5925.720
4862.187
February    -6139.6753   2386.993    -2.572    0.013   -1.09e+04

```

-1356.034					
January	-4482.1482	1405.485	-3.189	0.002	-7298.803
-1665.494					
July	2143.4046	1439.102	1.489	0.142	-740.619
5027.429					
June	-1495.5478	2404.016	-0.622	0.536	-6313.304
3322.209					
March	-1390.2979	2396.046	-0.580	0.564	-6192.081
3411.485					
May	2579.6164	2403.702	1.073	0.288	-2237.510
7396.743					
November	-3629.1572	2465.801	-1.472	0.147	-8570.733
1312.419					
October	-822.9414	1416.882	-0.581	0.564	-3662.436
2016.553					
September	-3427.1101	2435.594	-1.407	0.165	-8308.150
1453.929					
=====					
Omnibus:	11.113	Durbin-Watson:	1.527		
Prob(Omnibus):	0.004	Jarque-Bera (JB):	12.869		
Skew:	0.716	Prob(JB):	0.00161		
Kurtosis:	4.496	Cond. No.	3.06e+07		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.06e+07. This might indicate that there are strong multicollinearity or other numerical problems.

/tmp/ipykernel\_357/1113378029.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
civic['interest'].fillna(civic['interest'].mean(), inplace=True)
```

**Results and Observations** Based on the table above, our R-squared value is **0.807**, which is actually slightly better than all our above models from the previous questions (with model 2 being 0.806 and model 3 being 0.805). From the new variable we added named 'interest' (standing for the monthly interest rates on auto loans), we can observe that a 1% increase leads to a -327.098

decrease in Civic Sales. This is in line with the original thought I had when choosing this variable. With interest rates increasing, people who cannot afford cars are less likely to borrow money to buy cars since interest rates are higher.

Just an elaboration, but the dataset I had imported had some missing data for some months, but it did basically reflect all the years our original dataset had. To mitigate this, I took the mean of the interest rates and instead replaced the null values to get a better overall estimation of what the data would look like and for it to be integrable within our previous data's CSV.

sourcing: <https://fred.stlouisfed.org/categories/33058> <https://fred.stlouisfed.org/series/TERMCBAUTO48NS>

[ ]: